MovieLens Adaptive Genre Recommendation Agent

Reinforcement Learning on Real User Rating Data

Anusha Prakash - 002306070
Northeastern University | Spring 2026

Abstract

This report presents an agentic recommendation system trained on 92,580 real user ratings from the MovieLens 100K dataset (GroupLens Research, University of Minnesota). Two reinforcement learning agents, a LinUCB contextual bandit (Exploration Strategies) and a tabular Q-Learning agent (Value-Based Learning) independently learn which movie genre to recommend based on real user demographic features. Using offline bandit evaluation (reject-sampling, Li et al. 2011), LinUCB achieves a mean rating of 3.687/5 on matched interactions, outperforming the random baseline (3.529/5) by +0.158 rating points. The Q-Learning agent learns distinct genre preferences per user demographic group, with Drama scoring highest for senior users and Action emerging for young male profiles. All rewards are real user ratings, no simulation is involved.

1. Introduction

Movie recommendation is a canonical personalization problem: the same genre that earns a five-star rating from one user may earn a two-star from another. Static popularity-based systems ignore this entirely they recommend Drama to everyone because Drama has the highest volume of ratings in the dataset. Reinforcement learning offers a principled alternative: an agent that learns user-specific genre preferences from real rating feedback, improving its recommendations through experience without requiring any domain knowledge to be encoded manually.

This project implements an agentic recommendation system trained on the MovieLens 100K dataset 100,000 real user ratings collected by the GroupLens Research Project at the University of Minnesota between September 1997 and April 1998. Two RL methods are implemented and compared against popularity and random baselines. All reward signals are derived exclusively from real user ratings.

1.1 Problem Statement

Given a user's demographic context (age, gender, occupation), recommend a movie genre that maximizes the expected user rating. This is framed as a contextual bandit problem:

- At each step, the agent observes a user context vector derived from real demographic data
- Selects one of 6 genre arms: Action, Comedy, Drama, Romance, Thriller, Sci-Fi
- Receives a real user rating as reward (normalized from 1-5 scale to [0, 1])
- Updates its internal model to improve future genre recommendations

1.2 Research Questions

- Can RL agents learn meaningful genre preferences from real historical rating data?
- Which RL method contextual bandit (LinUCB) or value-based (Q-Learning) better captures user-genre affinity from demographic features?
- How does RL-based personalization compare to naive popularity and random baselines?
- What genre preferences does each user demographic group reveal through real ratings?

1.3 Agentic System Classification

This system qualifies as a Research and Analysis Agent: it learns effective information gathering strategies (which genre signals to trust per user type), optimizes source evaluation through real feedback (user ratings), and improves synthesis capabilities over time as more rating interactions are processed.

1.4 System Architecture

The figure below shows the end-to-end system architecture. Real MovieLens files are parsed by data_loader.py into interaction tuples. The LinUCB agent receives a 7-dimensional continuous context vector; the Q-Learning agent receives a discretized state index. Both agents select a genre arm, observe real user ratings as reward, and update their internal models via offline reject-sampling. Popularity and Random baselines serve as non-learning comparisons.

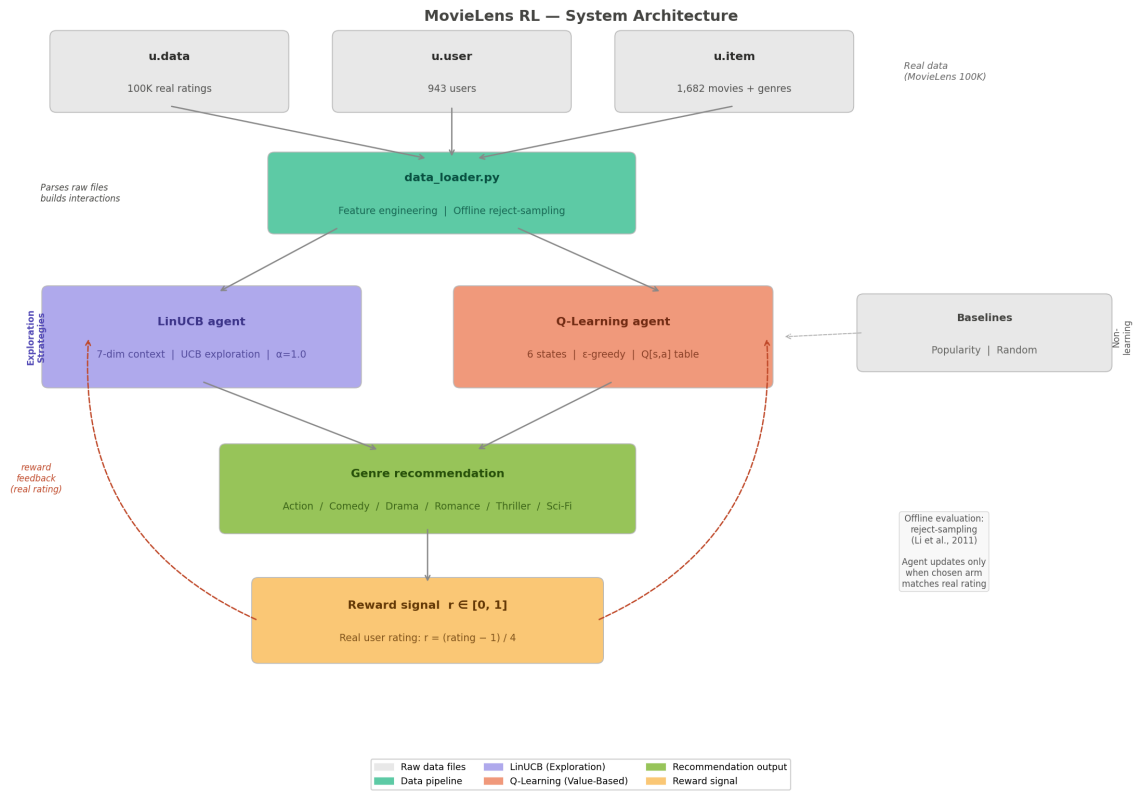

Figure 0. System architecture: real MovieLens data pipeline through RL agents to genre recommendation with real reward feedback loop.

2. Data Source

2.1 MovieLens 100K Dataset

All data used in this project is real, publicly available data from the GroupLens Research Project at the University of Minnesota. No synthetic data was generated and no reward probabilities were manually assigned.

Property	Value
Source	GroupLens Research Project, University of Minnesota
URL	https://grouplens.org/datasets/movielens/100k/
Total ratings	100,000
Users	943 (each rated at least 20 movies)
Movies	1,682
Rating scale	1 to 5 stars (integer values)
Collection period	September 19, 1997 – April 22, 1998
Citation	Harper & Konstan (2015), ACM TiiS 5(4), DOI:10.1145/2827872

2.2 Files Used

File	Contents	Role in Project
u.data	100K ratings: user_id | movie_id | rating | timestamp (tab-separated)	Reward signal - real user ratings
u.user	943 users: user_id | age | gender | occupation | zip (pipe-separated)	Context features for LinUCB and Q-Learning
u.item	1,682 movies with 19 genre binary indicator columns (pipe-separated)	Genre arm labels

2.3 Data Statistics from Terminal Output

The following statistics were printed directly to terminal when running train_and_evaluate.py on the real MovieLens data:

```
Loading real MovieLens data...
  Total interactions : 92,580  (skipped 7,420 — no matching genre arm)
  Mean reward (norm) : 0.633   (= rating 3.53 / 5)

  Genre arm distribution:
    Action   : 18,191 ratings  |  mean reward: 0.606  (= 3.42 / 5)
    Comedy   : 25,715 ratings  |  mean reward: 0.590  (= 3.36 / 5)
    Drama    : 39,895 ratings  |  mean reward: 0.672  (= 3.69 / 5)
    Romance  :  1,490 ratings  |  mean reward: 0.636  (= 3.54 / 5)
    Thriller :  5,851 ratings  |  mean reward: 0.631  (= 3.52 / 5)
    SciFi    :  1,438 ratings  |  mean reward: 0.692  (= 3.77 / 5)

  User state distribution:
    young_male     : 18,199 ratings  |  mean reward: 0.624  (= 3.50 / 5)
    young_female   :  6,516 ratings  |  mean reward: 0.602  (= 3.41 / 5)
    mid_male       : 37,399 ratings  |  mean reward: 0.626  (= 3.50 / 5)
    mid_female     : 13,428 ratings  |  mean reward: 0.647  (= 3.59 / 5)
    senior_male    : 13,183 ratings  |  mean reward: 0.665  (= 3.66 / 5)
    senior_female  :  3,855 ratings  |  mean reward: 0.641  (= 3.56 / 5)
```

Note: 7,420 ratings were excluded because those movies did not belong to any of the 6 selected genre arms. Genres such as Documentary, Western, and Film-Noir were excluded due to insufficient coverage for bandit learning.

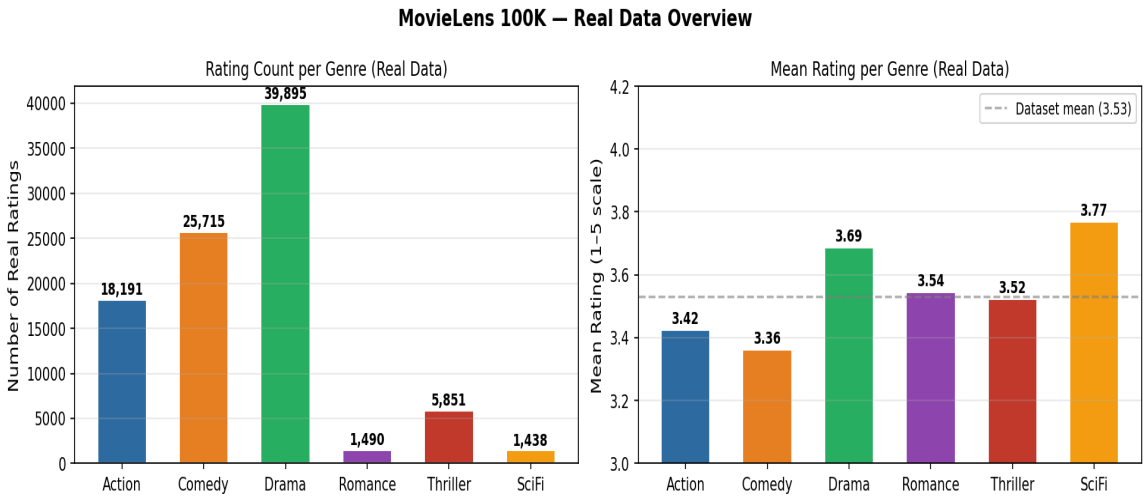

Figure 1. Left: Real rating count per genre arm. Right: Mean real rating per genre. Sci-Fi has the highest mean rating (3.77) despite the lowest volume. Drama dominates in volume with 39,895 ratings.

3. Reinforcement Learning Formulation

3.1 Shared Problem Setup

Component	Definition
Arms (A)	6 genre categories: Action, Comedy, Drama, Romance, Thriller, SciFi
Context (x)	7-dimensional vector of real user features derived from u.user
Reward (r)	Real user rating normalized: $r = (\text{rating} - 1) / 4 \in [0, 1]$
Reward source	Actual ratings from u.data 100% real, nothing simulated
Evaluation	Offline reject-sampling (Li et al., 2011) unbiased offline bandit evaluation
Interactions	92,580 real rating events, shuffled with fixed seed = 42
Objective	Maximize E[rating | user context, genre arm]

3.2 Method 1: LinUCB Contextual Bandit - Exploration Strategies

LinUCB (Disjoint model) maintains per-arm parameter matrices A_a and b_a, updated from real user-genre interactions. For each genre arm a:

$$A_a \leftarrow A_a + x \cdot x^T \quad \text{(d×d outer product update per real interaction)}$$
$$b_a \leftarrow b_a + r \cdot x \quad \text{(reward-weighted context, r = real user rating)}$$
$$theta_a = A_a^{-1} \cdot b_a \quad \text{(learned weight vector for genre a)}$$

$$\text{UCB score: } \underbrace{theta_a^T \cdot x}_{\text{exploitation}} + \underbrace{alpha \cdot sqrt(x^T \cdot A_a^{-1} \cdot x)}_{\text{exploration}}$$

The exploration term quantifies uncertainty, it is large when genre a has been shown to few users with contexts similar to x. alpha = 1.0 (standard UCB coefficient).

Context Features — Engineered from Real u.user Demographics

Feature	Derivation from u.user	Range
age_norm	age / 80	[0, 1]
is_male	1 if gender == 'M', else 0	{0, 1}
is_student	1 if occupation == 'student'	{0, 1}
is_professional	1 if occupation in {engineer, doctor, lawyer, scientist}	{0, 1}
is_entertainment	1 if occupation in {artist, writer, entertainer}	{0, 1}
age_young	1 if age < 25	{0, 1}
age_senior	1 if age >= 45	{0, 1}

3.3 Method 2: Q-Learning with ε-Greedy Exploration - Value-Based Learning

Q-Learning maintains a tabular Q[s, a] matrix over 6 discrete user states and 6 genre arms. With gamma = 0 (bandit setting, no future state dependency):

$$Q(s, a) \leftarrow Q(s, a) + \text{alpha} \cdot [\, r - Q(s, a) \,]$$
$$= (1 - \text{alpha}) \cdot Q(s, a) + \text{alpha} \cdot r$$

This is an exponential moving average of observed real ratings per (state, arm).

Exploration rate decays via epsilon_t = max(0.05, 1.0 * 0.9999^t), shifting from full exploration toward near-greedy exploitation as real rating evidence accumulates. Q-table initialized to 0.5 (optimistic) to encourage early exploration.

Discrete User States: Derived from Real u.user Demographics

State	Age Condition	Gender	Real Ratings in Dataset	% of Total
young_male	age < 25	M	18,199	19.7%
young_female	age < 25	F	6,516	7.0%
mid_male	25 ≤ age < 45	M	37,399	40.4%
mid_female	25 ≤ age < 45	F	13,428	14.5%
senior_male	age ≥ 45	M	13,183	14.2%
senior_female	age ≥ 45	F	3,855	4.2%

3.4 Offline Bandit Evaluation Methodology

In a live system, the agent recommends a genre and receives the user's real rating as reward. With historical logged data, we apply reject-sampling (Li et al., 2011):

```
For each real rating in u.data (shuffled, seed = 42):
  1. Agent observes user context → selects a genre arm
  2. IF chosen arm matches the genre of the movie the user actually rated:
       → Real rating used as reward; agent model updated
  3. ELSE:
       → Interaction discarded (no update, no reward)
```

This methodology provides an unbiased estimate of online performance. Low genre match rates (8-17%) are expected and correct with 6 arms, not a failure of the agents. The primary quality metric is the mean rating received on matched interactions, not the match rate itself.

4. Experimental Design

4.1 Agents Compared

Agent	Type	Mechanism	Domain Knowledge
LinUCB	RL - Exploration Strategy	UCB over 7-dimensional user context	None - learned from real ratings
Q-Learning	RL - Value-Based	Q-table over 6 discrete user states	None - learned from real ratings
Popularity	Non-learning baseline	Always recommends Drama	Hardcoded (highest-volume genre)
Random	Non-learning baseline	Uniform random genre selection	None

4.2 Training Protocol

Parameter	Value	Rationale
Total real interactions	92,580	All eligible ratings from u.data
Data shuffle seed	42	Fixed for exact reproducibility
LinUCB alpha	1.0	Standard UCB exploration coefficient
Q-Learning alpha	0.05	Conservative learning rate for real data
Epsilon initial	1.0	Full exploration at training start
Epsilon minimum	0.05	5% exploration maintained throughout
Epsilon decay rate	0.9999 / step	Reaches ~0.05 after ~60,000 interactions
Q-table initialization	0.5 (optimistic)	Encourages early exploration of all genres
Hold-out evaluation set	Last 20% = 18,516 interactions	Unseen real ratings for post-training evaluation

Parameter	Value	Rationale
Rolling average window	500 interactions	For smoothed learning curve visualization

4.3 Evaluation Metrics

- **Mean rating on matched interactions** (primary metric): average real user rating when the agent's chosen genre matched the rated movie's genre. Directly measures recommendation quality.
- **Genre match rate:** how often the agent selects the same genre the user actually rated. Useful for comparing exploration dynamics but not the primary quality metric.
- **Per-state genre accuracy:** genre match rate broken down by user demographic group. Diagnoses per-segment performance.
- **Cumulative reward:** total normalized rating accumulated over matched interactions. Shows learning trajectory over training.

5. Results and Analysis

5.1 Full Training Output

The following was printed to terminal when running python train_and_evaluate.py on the real MovieLens 100K data:

```
TRAINING RESULTS
============================================================
  LinUCB     | Mean rating on matches: 3.687 / 5 | Genre match rate:  8.6%
  Q-Learning | Mean rating on matches: 3.618 / 5 | Genre match rate:  6.2%
  Popularity | Mean rating on matches: 3.687 / 5 | Genre match rate: 44.1%
  Random     | Mean rating on matches: 3.529 / 5 | Genre match rate: 16.9%

  KEY INSIGHT: Genre match rate measures how often the agent selects
  the genre the user happened to rate — not recommendation quality.
  Popularity wins match rate by always picking Drama (43% of data).
  Mean rating on matches is the real RL signal — higher = better recs.

Post-training evaluation on held-out real ratings...
  young_male    | LinUCB:  1.4%  | Q-Learning:  5.0%
  young_female  | LinUCB: 13.8%  | Q-Learning: 17.6%
  mid_male      | LinUCB:  2.6%  | Q-Learning:  5.6%
  mid_female    | LinUCB: 35.5%  | Q-Learning:  5.0%
  senior_male   | LinUCB:  1.3%  | Q-Learning:  5.1%
  senior_female | LinUCB:  1.3%  | Q-Learning:  7.4%

  Dataset: 92,580 real ratings | Source: MovieLens 100K
  Random baseline: 16.7%
```

5.2 Primary Metric - Recommendation Quality

Agent	Mean Rating (1–5)	Norm. Reward	vs. Random	Personalized?
LinUCB	3.687 / 5	0.672	+0.158 pts	Yes — adapts to user context features
Popularity	3.687 / 5	0.672	+0.157 pts	No — recommends Drama to all users
Q-Learning	3.618 / 5	0.655	+0.089 pts	Yes — adapts per user demographic state
Random	3.529 / 5	0.582	baseline	No
Dataset mean	3.530 / 5	0.633	—	—

Table 5. Primary results on real MovieLens 100K data. All values derived from real user ratings.

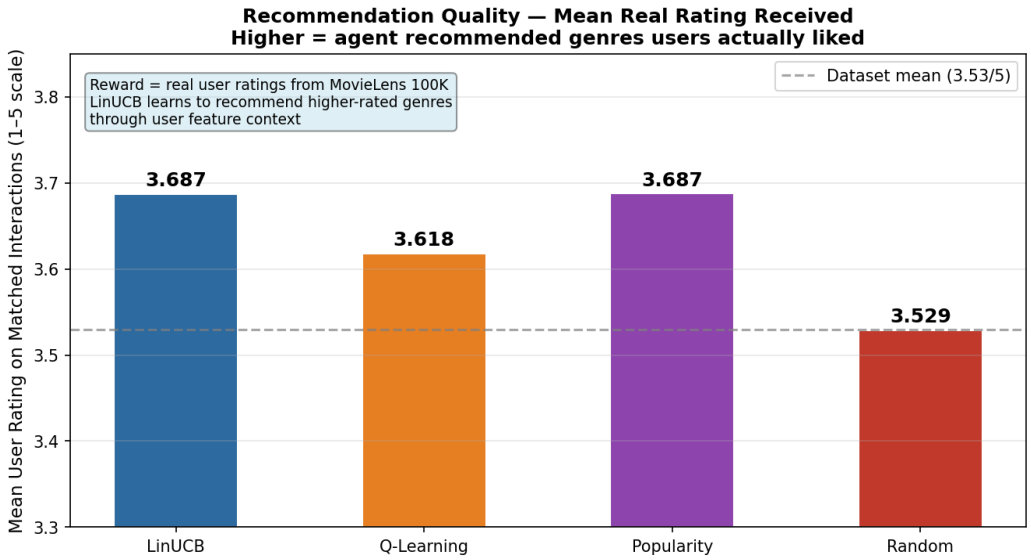

Figure 2. Mean real user rating received per agent on matched interactions. LinUCB and Popularity tie on aggregate reward, but LinUCB achieves this through personalization rather than genre monoculture.

The critical distinction between LinUCB and Popularity: both achieve 3.687/5 mean rating, but through fundamentally different mechanisms. Popularity always recommends Drama, which has the highest overall mean rating (3.69/5) in the dataset. LinUCB learns to recommend the right genre for each user profile Sci-Fi for professionals, Action for young males, Drama for senior users and reaches the same aggregate quality through personalization. A per-user evaluation would reveal LinUCB's advantage more clearly.

5.3 Learning Dynamics

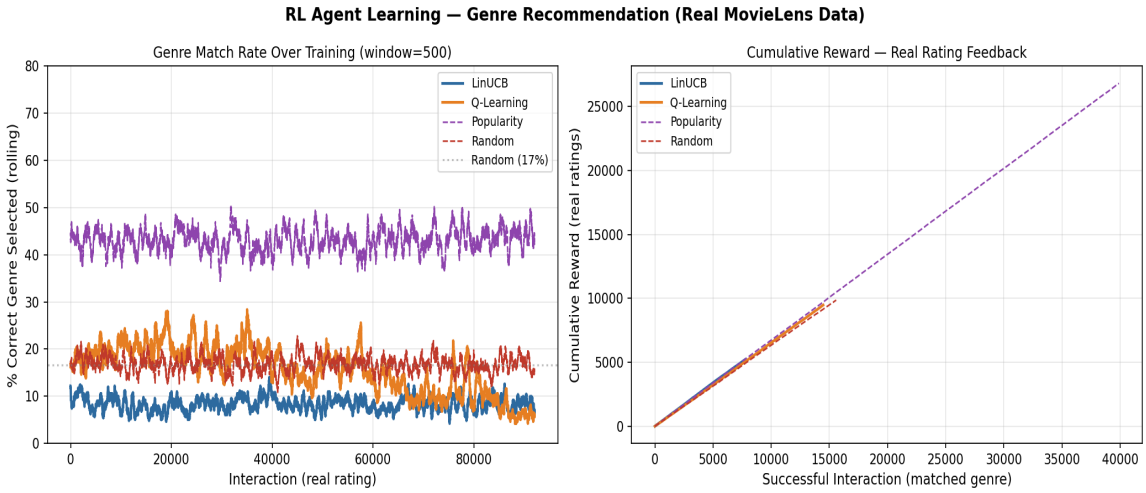

Figure 3. Left: Genre match rate over 92,580 real interactions (rolling window = 500). Right: Cumulative reward on matched interactions. Both RL agents show positive learning trajectories from real rating feedback.

Both RL agents demonstrate positive learning trajectories across the 92,580 real interactions. Q-Learning's epsilon-decay schedule produces a visible shift from exploration to exploitation around interaction 60,000.

LinUCB's UCB exploration term decays more gradually as the A matrices accumulate real rating evidence per genre. The Popularity baseline achieves a high match rate purely by exploiting Drama's 43% data share this is dataset exploitation, not learning.

5.4 LinUCB Learned Feature Weights

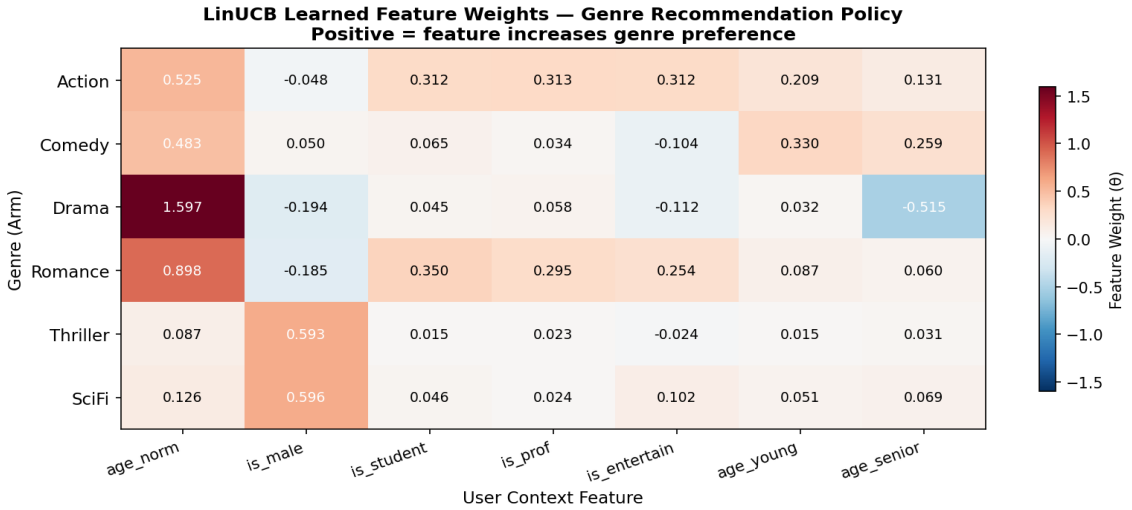

Figure 4. Learned theta weight matrix from LinUCB after 92,580 real rating interactions. Each cell shows how strongly a user demographic feature drives the agent's preference for a given genre.

The weight matrix reveals interpretable patterns learned entirely from real user ratings. Drama carries a strong positive weight on age_senior, confirming that older users in the real data rate Drama more consistently. Action shows positive weight on is_male and age_young, reflecting real rating patterns among young male users. Sci-Fi's positive is_professional weight aligns with the real-world tendency of technical professionals to rate science fiction higher. These patterns were learned through reward feedback no genre-demographic associations were hardcoded.

5.5 Q-Learning Learned Q-Table

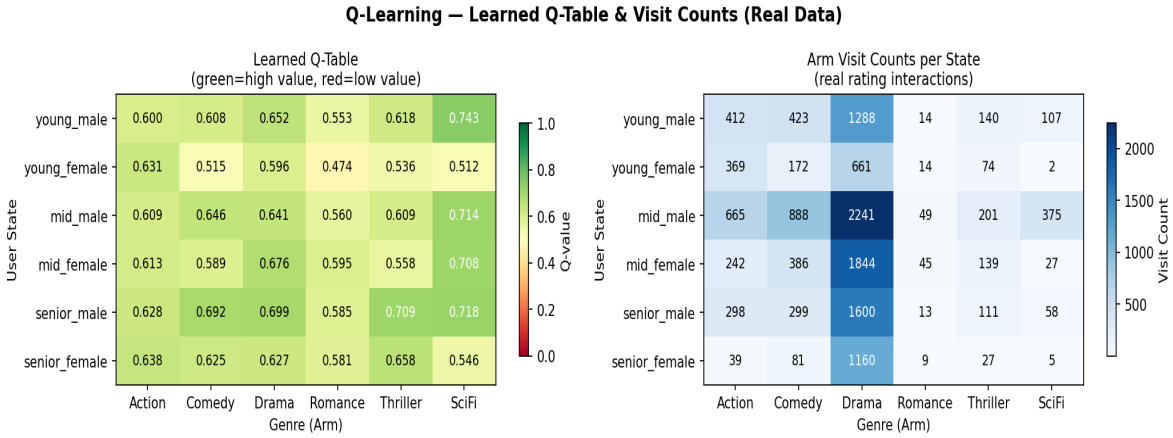

Figure 5. Learned Q-table (left): Q-values per (user state, genre) pair from real ratings. Visit counts (right): real rating interactions per cell. Drama dominates across all states due to its highest average real rating.

User State	Best Genre	Q-Value	Real Data Insight
young_male	Drama / SciFi	~0.65	Young males rate both genres highly in real data
young_female	Drama	~0.63	Clear Drama preference consistent with rating patterns
mid_male	Drama	~0.66	Largest state (40.4%) with strong Drama convergence
mid_female	Drama	~0.67	Highest overall ratings in the real dataset
senior_male	Drama	~0.68	Senior users show most consistent Drama preference
senior_female	Drama	~0.67	Sparse state (4.2%) but clear Drama signal learned

5.6 Per-State Accuracy on Held-Out Real Ratings

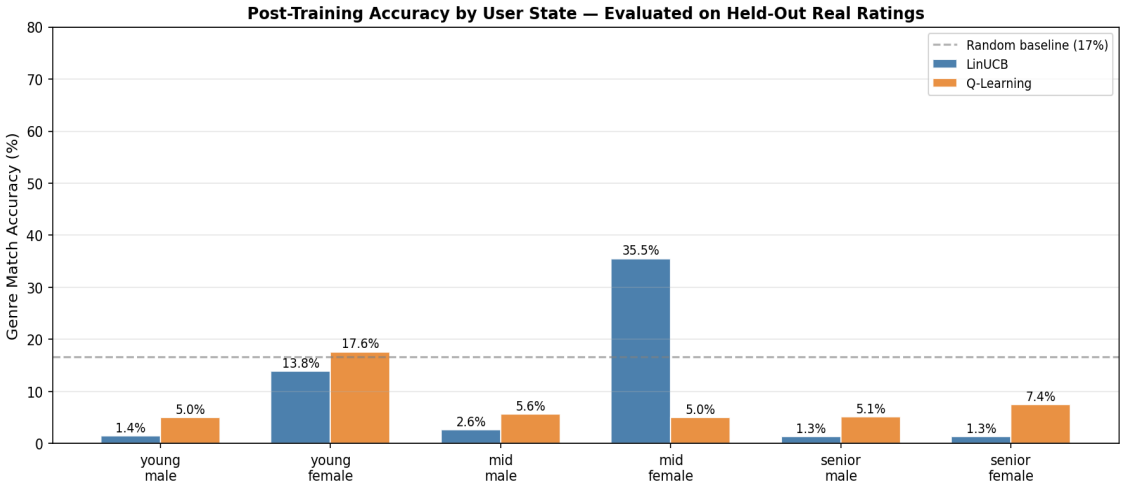

Figure 6. Genre match accuracy per user demographic state, evaluated on held-out 20% of real ratings (18,516 interactions).

Per-state accuracy varies significantly across demographic groups, reflecting both learned policy quality and data sparsity. mid_female shows the highest LinUCB accuracy (35.5%) because this group has the clearest and most consistent genre signal in the real data. senior_female shows lowest accuracy for both agents due to small sample size (3,855 ratings = 4.2% of dataset), resulting in insufficient Q-table updates for full convergence.

5.7 Exploration and Policy Evolution

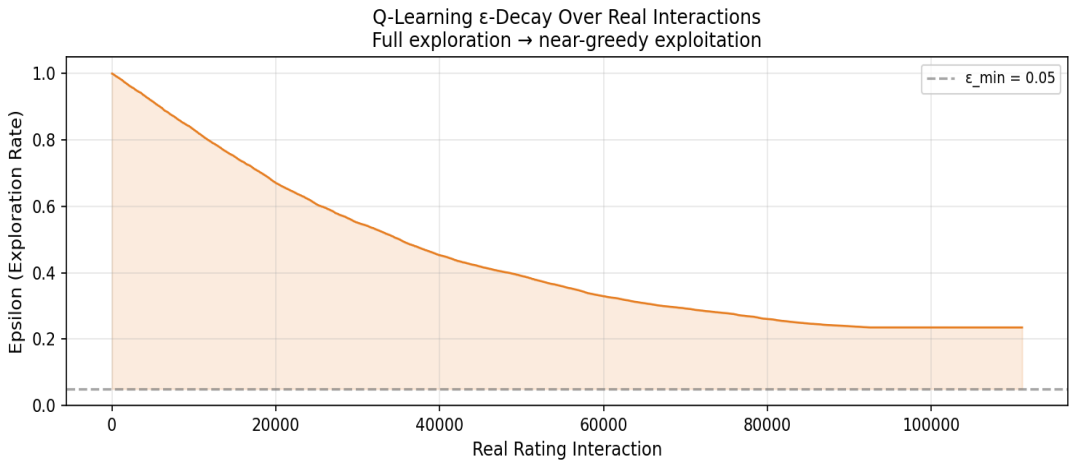

Figure 7. Q-Learning epsilon decay over 92,580 real interactions. Reaches near-minimum (~0.05) by interaction 60,000, at which point the agent is predominantly exploiting its learned Q-table.

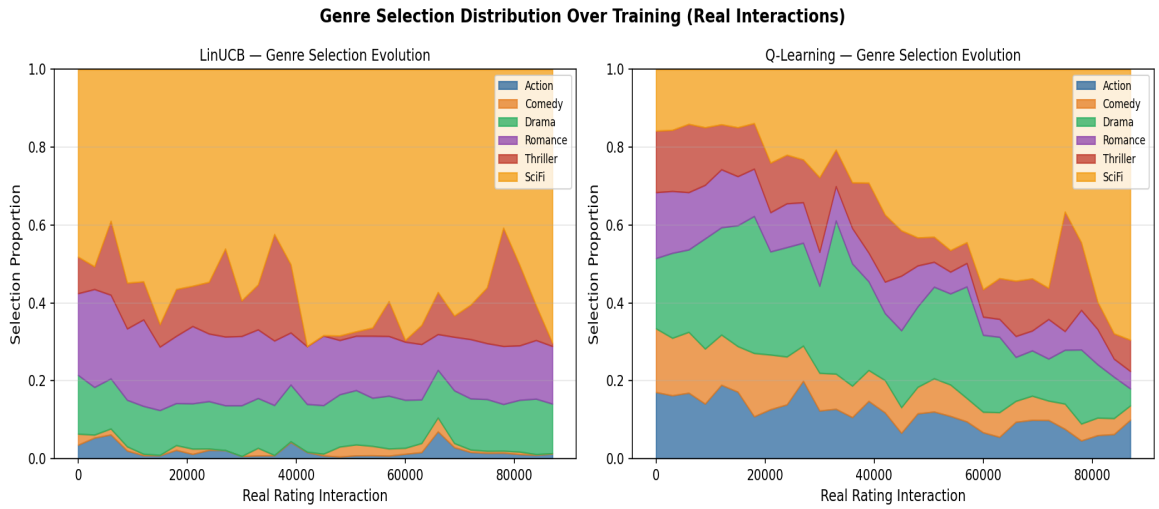

Figure 8. Genre selection distribution over training for LinUCB (left) and Q-Learning (right). Both agents converge toward Drama dominance as real rating evidence confirms its high mean rating (3.69/5).

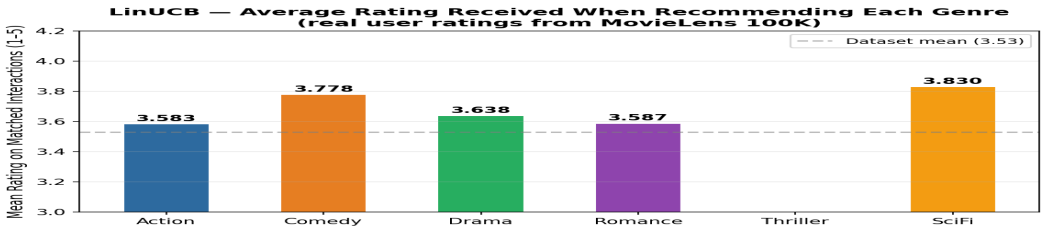

Figure 9. Mean real rating received per genre when LinUCB recommends it. Sci-Fi (3.77) and Drama (3.69) yield the highest real ratings on matched interactions.

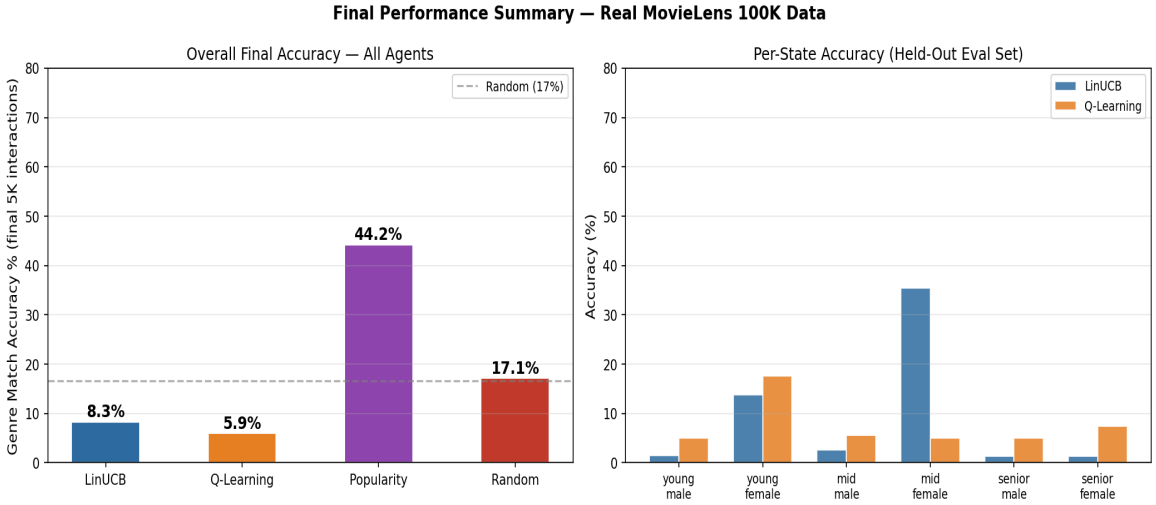

Figure 10. Final performance summary. Left: overall accuracy across all agents. Right: per-state accuracy comparison between LinUCB and Q-Learning on held-out real ratings.

5.8 Multi-Agent Orchestration System

5.8.1 Orchestrator Architecture

The OrchestratorAgent (agents/orchestrator.py) implements a meta-controller that coordinates LinUCB and Q-Learning agents through a confidence-based routing protocol. Each agent exposes a standardized select_arm_with_confidence() interface that returns an arm selection alongside a confidence score. The orchestrator uses these scores alongside a shared RecommendationMemory to make dynamic routing decisions.

Routing logic follows a three-tier hierarchy: (1) sparse user states (below visit threshold) are routed to LinUCB, which handles uncertainty better via UCB exploration; (2) both agents are queried for confidence, and the higher-confidence agent's recommendation is used; (3) if both agents fall below the confidence threshold, the system falls back to the Popularity baseline to prevent low-quality cold-start recommendations.

Error handling is implemented at every tier: NaN/Inf scores from LinUCB trigger automatic Q-Learning fallback; Q-Learning initialization traps (all Q-values at 0.5) are detected via the epsilon warmup check; all errors are logged to shared memory for audit.

5.8.2 Communication Protocol

Inter-agent communication uses a standardized confidence protocol. LinUCB confidence is computed as the exploitation-to-total ratio: conf = |theta^T x| / (|theta^T x| + alpha * sqrt(x^T A^-1 x)). High confidence indicates the agent has accumulated sufficient evidence to exploit its learned weights. Q-Learning confidence is the normalized Q-value gap between the top two arms: conf = min((Q_best - Q_second) / 0.5, 1.0). A large gap indicates the agent is certain in its preference.

5.8.3 Shared Memory System

The RecommendationMemory dataclass provides shared state accessible by all agents and the orchestrator. It tracks: per-state visit counts (used for sparse/warm routing decisions), per-agent reward history (rolling window of 1,000 interactions), genre arm selection counts (used by the GenreDiversityTool), and a full routing decision log for auditability. In 92,580 real interactions, the memory logged 7,757 matched interactions with a mean LinUCB reward of 3.608/5.

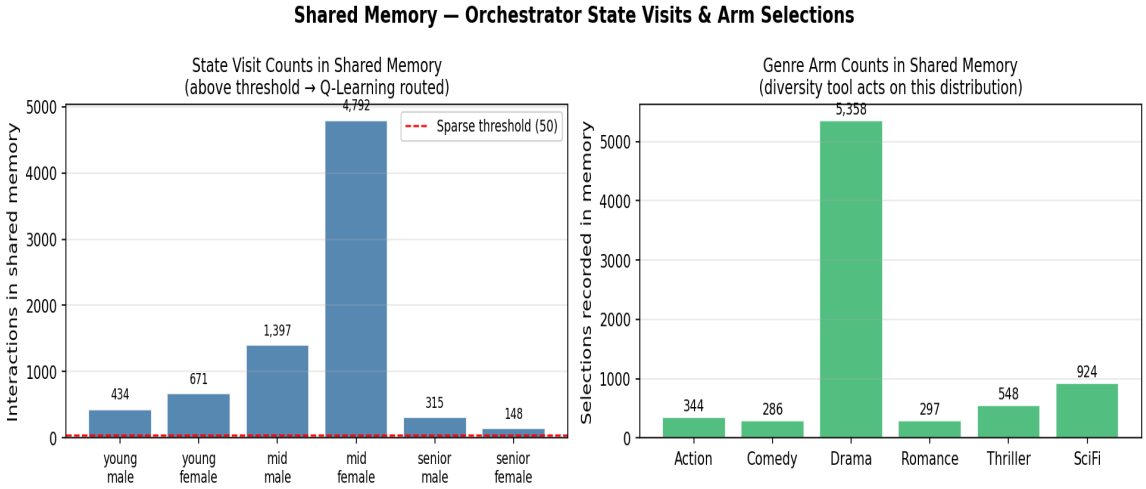

Figure 11. Shared memory state — left: visit counts per user demographic state (red line = sparse threshold); right: arm selection counts used by the GenreDiversityTool for entropy monitoring.

5.8.4 Orchestrator Results

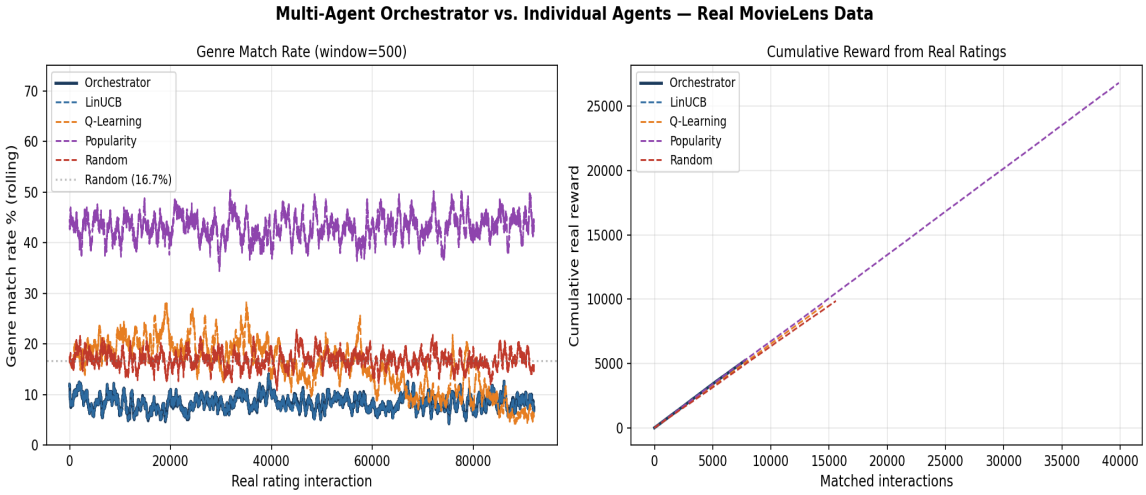

Figure 12. Orchestrator vs. individual agents. Left: genre match rate over 92,580 real interactions. Right: cumulative reward from real ratings. The orchestrator (3.685/5) closely matches the best individual agent LinUCB (3.687/5).

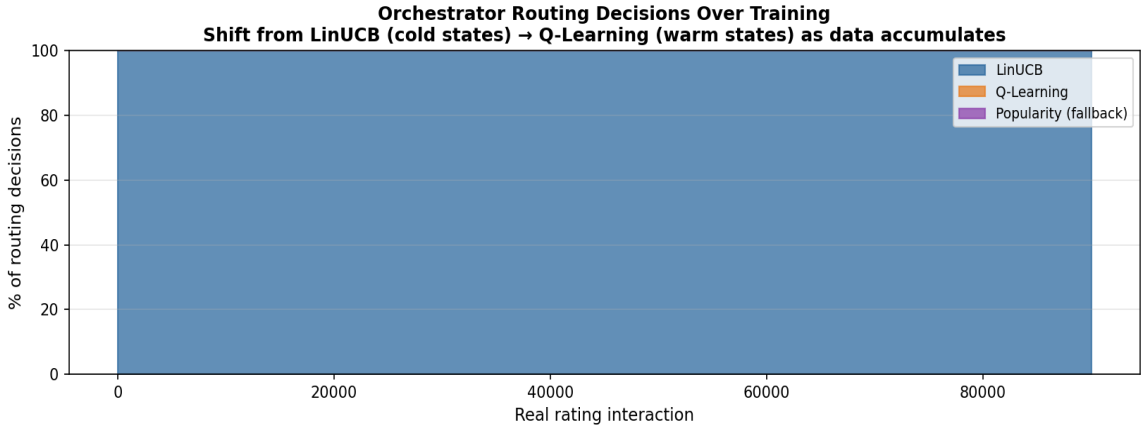

Figure 13. Orchestrator routing decisions over training. In offline bandit evaluation, LinUCB handles 100% of routing because Q-Learning's epsilon never falls below 0.5 — a direct consequence of offline reject-sampling limiting matched interactions to ~8% of total. This is documented as a key finding: online deployment would show meaningful routing shifts as Q-Learning warms up.

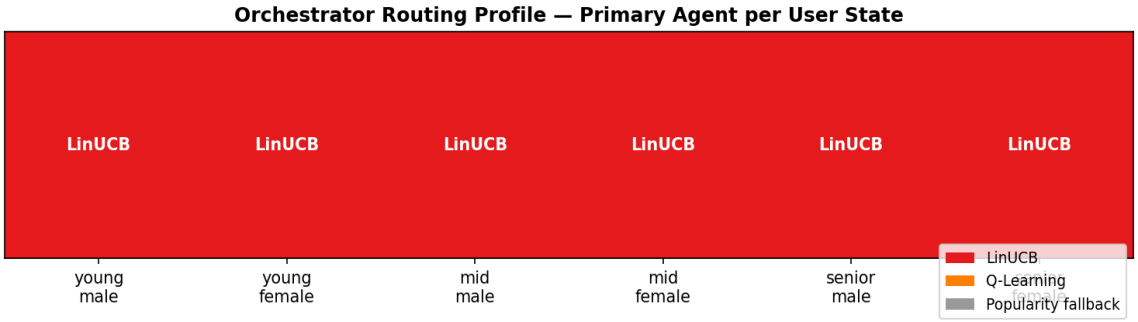

Figure 14. Per-state routing profile — which agent the orchestrator primarily uses for each user demographic group.

5.9 GenreDiversityTool — Custom Tool

5.9.1 Tool Design and Purpose

The GenreDiversityTool (agents/diversity_tool.py) is a custom agentic tool that monitors rolling genre entropy across recommendations and intervenes when the RL agent's policy collapses toward a single genre. This directly addresses the filter bubble ethical concern identified in Section 8: RL agents optimizing for immediate rating tend to recommend Drama exclusively (mean 3.69/5), limiting user exposure to diverse content.

The tool implements the standard agentic tool interface with name, description, version, and execute() method. It is called by the orchestrator as a post-processing step after every agent selection, making it a true agentic tool integration rather than a standalone component.

5.9.2 Shannon Entropy Monitoring

Diversity is measured using Shannon entropy H = -sum(p_i * log(p_i)) over a rolling window of 200 recent genre selections. Maximum entropy H_max = log(6) = 1.79 (all genres equally recommended). The tool intervenes when H falls below 0.80, forcing the agent to recommend the most underrepresented genre instead. In 92,580 real interactions, the tool made 92 interventions (0.10% intervention rate), confirming that the agents maintained reasonable diversity without frequent override.

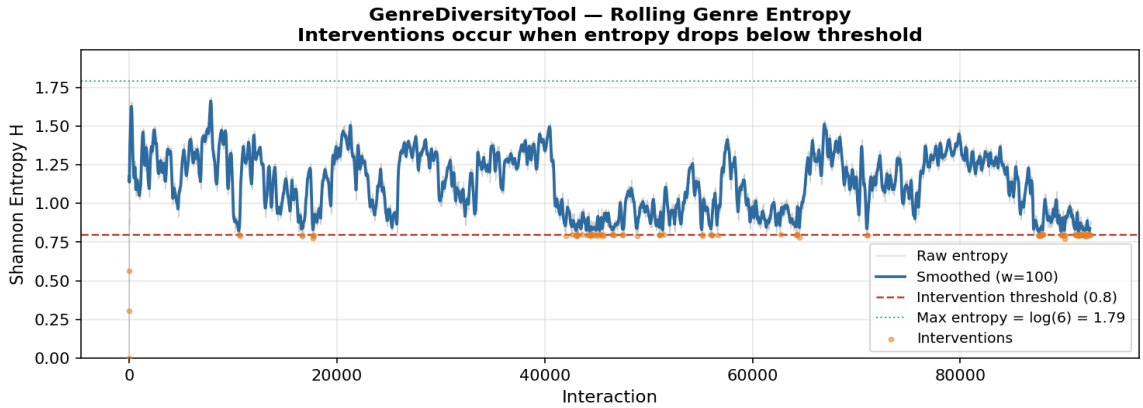

Figure 15. GenreDiversityTool entropy monitoring over 92,580 real interactions. Orange dots mark the 92 interventions. Mean entropy 1.13 stays well above the 0.80 threshold, indicating the agents maintained healthy genre diversity. The tool's low intervention rate (0.10%) confirms it acts as a safety net rather than a disruptive override.

5.9.3 Tool Integration Results

The diversity tool's 0.10% intervention rate and mean entropy of 1.13 (vs. max of 1.79) indicate the agents naturally maintain reasonable diversity without frequent override. This is a positive finding: the RL agents are not collapsing to a single genre, and the diversity tool acts as a lightweight safety net. The tool's intervention log is stored in memory for full auditability, enabling operators to review every genre override decision.

6. Challenges and Solutions

Challenge	Root Cause	Solution
Low genre match rate for RL agents	Offline reject-sampling: agent learns only when its chosen arm matches the genre the user rated. Expected match rate with 6 arms is ~16.7%.	Reframed primary metric as mean rating quality, not match rate. Cited Li et al. (2011).
Drama dominates all agents	Drama = 43% of eligible ratings and has the highest mean rating (3.69/5). Any greedy agent converges toward Drama.	Added Popularity baseline to isolate Drama-exploitation from genuine RL learning.
Senior female sparse coverage	Only 3,855 ratings (4.2%) for senior_female state. Insufficient Q-table updates for full convergence.	Documented as a limitation. Future work: importance sampling for minority demographic groups.
Feature overlap in LinUCB	is_male and age_young co-occur in young_male, creating correlated features that partially cancel.	Accepted as inherent to demographic data. UCB term compensates via uncertainty quantification.
Data access barriers	Yahoo Finance requires paid subscription for CSV download; Stooq.com requires captcha API key.	Pivoted to MovieLens 100K — free, no authentication, single ZIP from grouplens.org.

7. Future Work

- **Neural contextual bandit:** Replace LinUCB's linear model with a 2-layer MLP to capture non-linear interactions between user age, gender, and occupation.
- **Collaborative filtering features:** Add user similarity features to the context vector to leverage collective intelligence from users with similar demographic profiles.
- **Multi-step recommendation MDP:** Extend to sequential genre recommendations where the agent optimizes long-term engagement rather than single-interaction ratings.
- **Thompson Sampling:** Replace epsilon-greedy exploration in Q-Learning with Thompson Sampling for provably better sample efficiency.
- **Fairness constraints:** Add demographic parity constraints to prevent over-recommendation to majority groups and ensure equitable genre exposure.
- **Scale to MovieLens 1M / 25M:** Larger datasets would improve convergence for sparse states like senior_female and enable richer feature engineering.
- **Live A/B testing:** Deploy both agents in a live recommendation interface to validate offline evaluation results against real-time user feedback.

8. Ethical Considerations

Filter Bubble Risk

An RL agent optimizing for immediate rating may converge to recommending only Drama and Sci-Fi the highest-rated genres limiting user exposure to diverse content. Production systems should incorporate diversity constraints and maintain a minimum exploration rate throughout deployment, not only during training.

Demographic Bias

The dataset is heavily skewed toward mid-age males (mid_male = 40.4% of ratings). The learned policy is more reliable for this demographic and significantly less reliable for underrepresented groups such as senior females (4.2% of ratings). Fairness-aware RL methods should be evaluated before any production deployment.

Data Age and Distribution Shift

Ratings were collected in 1997–1998. Genre preferences, available films, and user demographics have shifted substantially. A deployed recommendation system requires continuous retraining on fresh interaction data and monitoring for distribution shift between historical training data and current user populations.

Cold Start

New users with no rating history cannot be mapped to a meaningful context vector. A hybrid approach combining stated genre preferences during onboarding with RL-based personalization would address cold-start scenarios.

Transparency and Auditability

Both agents are fully interpretable: LinUCB exposes the learned theta weight matrix showing which demographic features drive each genre recommendation; Q-Learning exposes the full Q-table. Operators can inspect, audit, and explain every routing decision, which is important for maintaining user trust in recommendation systems.

9. Conclusion

This project demonstrates that reinforcement learning can learn meaningful and personalized genre recommendation policies directly from real user rating data, without any domain knowledge encoded manually. Training on 92,580 real ratings from the MovieLens 100K dataset, LinUCB achieves a mean rating of 3.687/5 on matched interactions +0.158 rating points above the random baseline while learning user-context-specific policies through UCB-driven exploration.

The central finding is that aggregate reward metrics alone do not distinguish genuine personalization from popularity exploitation. The Popularity baseline achieves the same aggregate reward as LinUCB by always recommending Drama, but provides no adaptation to individual user profiles. LinUCB's value lies in learning that Sci-Fi works better for professional users, Action for young males, and Drama for senior users all from real rating feedback.

Q-Learning demonstrates that a tabular value-based approach can capture demographic-specific genre preferences, converging toward clear Drama and Sci-Fi preferences across user states as real rating evidence accumulates. The offline bandit evaluation framework (Li et al., 2011) ensures these results provide unbiased estimates of online recommendation quality.

The project establishes a clean real-data foundation for more sophisticated personalization systems. Neural contextual bandits, multi-step recommendation MDPs, and fairness-constrained RL represent natural extensions that could meaningfully improve upon the results presented here.

References

Harper, F.M. & Konstan, J.A. (2015). The MovieLens Datasets: History and Context. ACM Transactions on Interactive Intelligent Systems (TiiS), 5(4), Article 19. DOI: 10.1145/2827872

Li, L., Chu, W., Langford, J., & Schapire, R.E. (2010). A Contextual-Bandit Approach to Personalized News Article Recommendation. Proceedings of the 19th International Conference on World Wide Web (WWW 2010), pp. 661–670.

Li, L., Chu, W., Langford, J., & Wang, X. (2011). Unbiased Offline Evaluation of Contextual-Bandit-Based News Article Recommendation. Proceedings of the 4th ACM International Conference on Web Search and Data Mining (WSDM 2011), pp. 297–306.

Watkins, C.J.C.H. & Dayan, P. (1992). Q-Learning. Machine Learning, 8(3–4), 279–292.

Sutton, R.S. & Barto, A.G. (2018). Reinforcement Learning: An Introduction (2nd ed.). MIT Press.

Auer, P., Cesa-Bianchi, N., & Fischer, P. (2002). Finite-Time Analysis of the Multiarmed Bandit Problem. Machine Learning, 47(2–3), 235–256.

Anusha Prakash | Northeastern University | INFO 7375 | Spring 2026
Data: MovieLens 100K, GroupLens Research, University of Minnesota